

STAT 206 Homework 4_Lihua_Xu

Due Monday, October 30, 5:00 PM

General instructions for homework: Homework must be completed as an R Markdown file. Be sure to include your name in the file. Give the commands to answer each question in its own code block, which will also produce plots that will be automatically embedded in the output file. Each answer must be supported by written statements as well as any code used. (Examining your various objects in the “Environment” section of RStudio is insufficient – you must use scripted commands.)

In lecture, we fit a gamma distribution to the weight of cat’s hearts. We did this by adjusting the parameters so that the theoretical values of the mean and variance matched the observed, sample mean and variance. Since the mean and variance are the first two moments of the distribution, this is an example of the method of moments for estimation.

The method of moments gives a point estimate $\hat{\theta}$ of the parameters θ . To use a point estimate, we need to know how precise it is, i.e., how different it would be if we repeated the experiment with new data from the same population. We often measure imprecision by the standard error, which is the standard deviation of the point estimates $\hat{\theta}$. (You saw the standard error of the mean in your introductory statistics classes, but we are not computing the standard error of the mean here.)

If we actually did the experiment many times, getting many values of $\hat{\theta}$, we could take their standard deviation as the standard error. With only one data set, we need to do something else. There is usually no simple formula for standard errors of most estimates, the way there is for the standard error of the mean. Instead, we will see how to approximate the standard error of for our estimate of the gamma distribution computationally.

We can draw random values from a gamma distribution using the `rgamma()` function. For example, `rgamma(n=35,shape=0.57,scale=15)` would generate a vector of 35 random values, drawn from the gamma distribution with “shape” parameter $a = 0.57$ and “scale” $s = 15$. By applying the estimator to random samples drawn from the distribution, we can see how much the estimates will change purely due to noise.

Part I - Estimates and standard errors

1. Write a function, `gamma.est`, which takes as input a vector of data values, and returns a vector containing the two estimated parameters of the gamma distribution, with components named `shape` and `scale` as appropriate.

```
gamma.est <- function(vector_data)
{mean_value <- mean(vector_data)
  variance_value <- var(vector_data)
  scale <- variance_value/mean_value
  shape <- mean_value^2/variance_value
  return(c(shape=shape,scale=scale))}
```

2. Verify that your function implements the appropriate formulas by showing that it matches the results from lecture for the cat heart data.

```
##Loading data for the cat
library(MASS)
data(cats)

##The code for obatining shape and scale from class
gamma.est_MM <- function(x) {
m <- mean(x); v <- var(x)
```

```
return(c(shape=m^2/v, scale=v/m))
}
gamma.est_MM(cats$Hwt)
```

```
##      shape      scale
## 19.0653121  0.5575862
##The results from my code.
gamma.est(cats$Hwt)
```

```
##      shape      scale
## 19.0653121  0.5575862
##My result matches the results from the calss.
```

3. Generate a vector containing ten thousand random values from the gamma distribution with $a = 19$ and $s = 0.56$. What are the theoretical values of the mean and of the variance? What are their sample values?

```
random.gamma <- rgamma(n=10000,shape=19,scale=0.56)
mean_random_gamma <- c(mean_random_gamma=mean(random.gamma))
mean_random_gamma
```

```
## mean_random_gamma
##      10.6289
```

```
var_random_gamma <- c(var_random_gamma=var(random.gamma))
var_random_gamma
```

```
## var_random_gamma
##      5.917286
```

##The theoretical values of mean and of the variance are 10.61778 and 6.114769 respectively.

```
mean_sample <- c(mean_catHwt=mean(cats$Hwt))
mean_sample
```

```
## mean_catHwt
##      10.63056
```

```
var_sample <- c(var_catHwt=var(cats$Hwt))
var_sample
```


```
## var_catHwt
##      5.927451
```

##The sample values of mean and of the variance are 10.63056 and 5.927451 respectively.

4. Plot the histogram of the random values, and add the curve of the theoretical probability density function.

```
hist(random.gamma, probability = TRUE)
curve(dgamma(x, shape = 19, scale = 0.56), add = TRUE)
```

Histogram of random.gamma



5. Apply your `gamma.est` function to your random sample. Report the estimated parameters and how far they are from the true values.

```
gamma.est(random.gamma)

##      shape      scale
## 19.0921319  0.5567164

##The "shape" and "scale" value for random.gamma data are 19.1259659 and 0.5573949 respectively.
##The true value for "shape" and "scale" are 19.0653121 and 0.5575862 respectively.
##Compare the two set of values, they didn't give too much difference.
##The scale values are kind of same and the shape values only have a ~0.06 difference.
```

6. Write a function, `gamma.est.se`, to calculate the standard error of your estimates of the gamma parameters, on simulated data drawn from the gamma distribution. It should take the following arguments: true shape parameter `shape` (or `a`), true scale parameter `scale` (or `s`), size of each sample `n`, and number of repetitions at that sample size `B`. It should return two standard errors, one for the shape parameter `a` and one for the scale parameter `s`. (These can be either in a vector or in a list, but should be named clearly.) It should call a function `gamma.est.sim` which takes the same arguments as `gamma.est.se`, and returns an array with two rows and `B` columns, one row holding shape estimates and the other row scale estimates. Your `gamma.est.se` function should not, itself, estimate any parameters or generate any random values

```
##Writing function gamma.est.se
gamma.est.se <- function(a,s,n,B)
{shape <- sd(gamma.est.sim(a,s,n,B)[1,])
scale <- sd(gamma.est.sim(a,s,n,B)[2,])
return(c(standard_deviation_shape=shape,standard_deviation_scale=scale))}
```

Part II - Testing with a *stub*

To check that `gamma.est.se` works properly, write a *stub* or *dummy* version of `gamma.est.sim`, which takes the correct arguments and returns an array of the proper size, but whose entries are fixed so that it's easy for us to calculate what `gamma.est.se` ought to do.

7. Write `gamma.est.sim` so that the entries in the first row of the returned array alternate between `shape` and `shape+1`, and those in the second row alternate between `scale` and `scale+n`. For example `gamma.est.sim(2,1,10,10)` should return (row names are optional)

```
##           [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10]
## shapes      2    3    2    3    2    3    2    3    2    3
## scales      1   11    1   11    1   11    1   11    1   11
```

and `gamma.est.sim(2,8,5,7)` should return

```
##           [,1] [,2] [,3] [,4] [,5] [,6] [,7]
## [1,]      2    3    2    3    2    3    2
## [2,]      8   13    8   13    8   13    8
```

```
##Writing function gamma.est.sim
gamma.est.sim <- function(a,s,n,B)
{shape.vector <- c()
 scale.vector <- c()
 for (j in 1:B)
 {if (j%%2!=0)
 {shape.vector <- c(shape.vector,a)
  scale.vector <- c(scale.vector,s)}
 else
 {shape.vector <- c(shape.vector,a+1)
  scale.vector <- c(scale.vector,s+n)}
 }
 return(rbind(shape.vector,scale.vector))}
```

8. Calculate the standard deviations of each *row* in the two arrays above.

```
apply(gamma.est.sim(2,1,10,10), 1, sd)
```

```
## shape.vector scale.vector
##    0.5270463    5.2704628
```

```
apply(gamma.est.sim(2,8,5,7), 1, sd)
```

```
## shape.vector scale.vector
##    0.5345225    2.6726124
```

9. Run your `gamma.est.se`, with this version of `gamma.est.sim`. Do its standard errors match the standard deviations you just calculated? Should they?

```
gamma.est.se(2,1,10,10)
```

```
## standard_deviation_shape standard_deviation_scale
##                0.5270463                5.2704628
```

```
gamma.est.se(2,8,5,7)
```

```
## standard_deviation_shape standard_deviation_scale
##                0.5345225                2.6726124
```

```
##The standard errors matches the standard deviations I just calculated.  
##Yes, they should.
```

Part III - Replacing the stub

10. Write the actual `gamma.est.sim`. Each of the B columns in its output should be the result of applying `gamma.est` to a vector of `n` random numbers generated by a different call to `rgamma`, all with the same shape and scale parameters.

```
gamma.est.sim <- function(a,s,n,B) {  
  return(replicate(B, gamma.est(rgamma(n, shape = a, scale = s))))  
}
```

11. Run `gamma.est.se`, calling your new `gamma.est.sim`, with `shape=2`, `scale=1`, `n=10` and `B=1e5`. Check that the standard error for `shape` is approximately 1.6 and that for `scale` approximately 0.54. Explain why your answers are not exactly 1.6 and 0.54.

```
gamma.est.se(2,1,10,1e5)
```

```
## standard_deviation_shape standard_deviation_scale  
##                1.6214402                0.5432306  
##My answer for the standard error of `shape` is approximately 1.6  
##and that for `scale` approximately 0.54.
```